

The Matchmakers: Data Analysis Report

D'Azevedo, Gloria
gad87@cornell.edu

Yadav, Pihu
py82@cornell.edu

December 5, 2016

1 Executive Summary

The goal of this analysis is to create a model that identifies whether or not a male and a female are a match in the speed dating context. As a side effect of developing this model, important traits in a partner for both male and female viewpoints are identified so that effective onboarding surveys for speed dating events or online dating will have high response and high accuracy rates. The data has several types of formats including ordinal variables, open ended text variables, and binary variables. The 3 types of models that are tested are a nonregularized linear model, a logistic model, and a regularized model using lasso. Some analysis was also done using the unsupervised method K-Nearest-Neighbors but the misclassification rates were very high so our current implementation and predictors used is not the best for this problem. In the future we hope to incorporate other variables indicative of a “match” such as one person calling the other or the two people going on a date. However, these fields have low response rates and there is high bias so they may not be completely reliable but still give good insights about the data. In future analyses we hope to utilize resampling techniques such as bootstrapping and cross validation for some of our model implementations to get more robust estimates of model coefficients.

2 Introduction

The main goal of this analysis is to develop a model that predicts whether or not two people (a male and a female) will be attracted to one another. In addition, we also want to find strong predictors of compatibility. These predictors can be traits that a person has, activities they enjoy, or how important is religion to the participant and their partner. Some obvious applications for this model will be organizing future speed dating events as well as online surveys that online dating websites or mobile applications use to determine potential matches for a participant. The ideal model should work on both genders, and the model should be sparse or relatively small so that the surveys and forms (where we get data) will be shorter so people are more likely to complete the whole thing truthfully when prompted.

The data used in the analysis is the Kaggle Speed Dating Experiment data set which was obtained from speed dating events conducted by Columbia Business School professors. (<https://www.kaggle.com/annavictoria/speed-dating-experiment>). There are 21 different speed dating events (waves) that occur between October 2002 and April 2004, each of which have between 10 and 45 participants, and the number of men and women are roughly

equal in each. Each participant is asked to fill out a total of 4 surveys in addition to questionnaires about their partners during the events. The first survey occurs when they sign up, the second is halfway through the speed dating event, the third is the first follow-up the day after the speed dating event so that they will get a list of their matches, and the fourth is the second follow-up after the speed dating event which is sent 3-4 weeks after they had been sent their matches. However, it is important to note that the response rates for each of the 4 surveys are not 100% and the missing data poses a huge problem for some of the algorithms.

The types of questions asked on the survey are mostly ordinal or categorical variables with a couple of open-ended text questions. Some of the open ended questions have a low completion response rate such as the age field and also the school that you completed your undergraduate degree. These fields therefore cannot be reliably used in the analysis because there could be inherent biases for why the participants did not respond. In the future, surveys can still ask these types of questions as a categorical or multiple-choice variable by giving participants a few age ranges that they can choose from so they would not be explicitly stating their age (a potential reason for nonresponse).

3 Data Cleaning Remarks

During the following report we will take on the convention that a “participant” is the person filling out the survey and the “partner” is the person they interact with during that particular round of the speed dating wave. When scraping data and analyzing the coefficients of the predictors in the model, we denote the “participant” as the female and the “partner” as the male, so that the data is not double counted. All females have interacted with each male once and vice versa so this is not a problem. There is only one incident of a person not completing her initial surveys.

For many questions in all the surveys, participants are asked in the surveys to divide up 100 points to 6 attributes by importance. The attributes are attractiveness, sincerity, intelligence, fun personality, ambition, and shared interests. However, there are several questions where the shared interests are either not recorded in our data or the survey forgot that attribute when asking the participant. It may be possible to impute that value for some participants if we assume that the question was asked on the survey but the data is not recorded by subtracting the sum of the points for the other attributes from 100 to get the points for the shared interests attribute as long as there is no non-response in the other questions.

While investigating trends in categorical data such as the “field_cd” column (corresponds to a numerical code for the field of study that each person has) or “race” (integer corresponding to a type of race), we noted that there were some NA values so we reassigned those NA values to be that predictor’s “Other” category. Even though we would be over-counting the trends in the “Other” category, we hope that the majority of participants had accurately written down their race or what field they classify themselves in. If the data field was an ordinal variable such as “tv” (measures the interest that a person has in watching TV on a scale from 1 to 10) we can reassign NA values to be 0, indicating that they had no interest in that activity and that’s why the participants left them blank. Obviously, imputing too much data affects the results significantly, so we only do that when trying to implement the KNN algorithm.

4 Initial Data Analysis and Methods

In the initial sign-up survey, participants are asked how important each of 17 different activities are to them. The activities include yoga, reading, and watching sports, and there may be high correlation between the interest level in some of the activities. For example, a person who likes theater may also really enjoy watching movies. These values are integer values from 0 to 10 and a higher value implies higher interest in the activity.

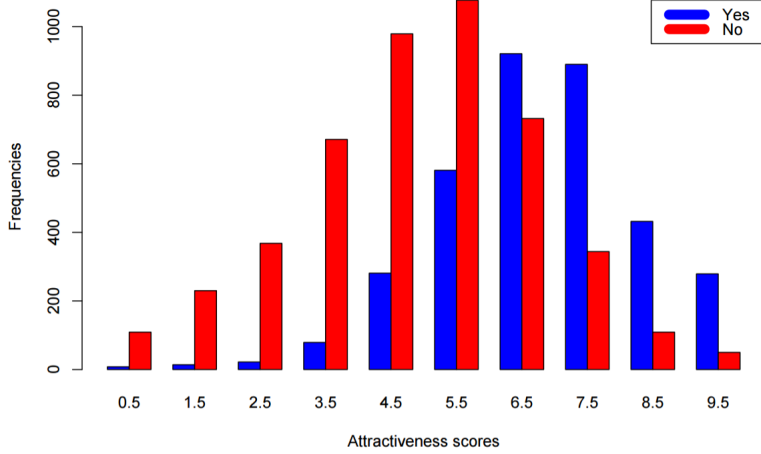
We tried fitting a linear model predicting match with all of the activities for both men and women (total of 34 variables) and taking only the significant variables from those to get a model. In addition, we try using a forward/backward selection algorithm using AIC to find an optimal model using the activities. Unfortunately we did not see that the same activities for both men and women were chosen so in the implementation we would have to have different onboarding surveys for men and women or we would include the nonsignificant variables on both surveys for the sake of consistency. We also considered using a vector of absolute differences ($\|\cdot\|_1$) or the sum of squared differences ($\|\cdot\|_2^2$) between the interest level in an activity between two people.

In one analysis, we identified how men value different traits that they want in a partner. Initially we computed the summary statistics for all the entries where men said 'yes' to their partner. This helped in getting a broad overview of the problem as it gave us the mean, median, maximum, minimum and some quantile values for the parameters we are interested in such as the scores given by the male participant for attributes such as attractiveness, sincerity, intelligence etc. to the partners they said 'yes' to. This helped us in getting an overall idea of what traits men value highly as high average or median scores in a particular category generally meant that men valued that trait. However, comparing summary values is not an adequate indicator of the significance of certain predictors, so we developed a more detailed model.

Table 1: Summary table for key statistics

Statistic	N	Mean	St. Dev.	Min	Max
samerace	3,518	0.409	0.492	0	1
int_corr	3,518	0.199	0.300	-0.830	0.910
attr	3,507	7.285	1.533	1.000	10.000
sinc	3,480	7.596	1.514	0.000	10.000
intel	3,474	7.757	1.330	3.000	10.000
fun	3,449	7.334	1.512	0.000	10.000
amb	3,271	7.159	1.592	0.000	10.000
shar	3,167	6.463	1.841	0.000	10.000

Figure 1: Histogram of Attractiveness Scores



5 Model Development

5.1 Logistic Regression

We want to identify how men value different traits that they want in a partner (attractiveness, sincerity, intelligence, being fun, ambition and shared interests). This is part of our attempts to understand gender differences in what participants look for in a partner. We also include the parameter “samerace” (which identifies whether the partner is of the same race) in the analysis.

Since the decision variable (of saying ‘yes’ or ‘no’ to a partner) is boolean we have chosen to use logistic regression for the analysis. The variables used include the scores given to partners (for attractiveness, sincerity, intelligence etc.) and whether or not the partner and the participant are of the same race. The results of this preliminary analysis suggested that attractiveness, sincerity, being fun, ambition and having shared interests are significant in a man’s decision and intelligence and being of the same race are not significant. Of these significant parameters, attractiveness, being fun and having shared interests are positively correlated with the final decision, but sincerity and ambition are negatively correlated with the male’s decision. In this analysis all the values with missing data have been removed from the data when computing the coefficients for the logistic regression, as it would be misleading to impute the values of scores given to partners or assign the average value when the answer has been left blank by the participant. We do not expect this to create problems in our analysis since more than 75% of the data is still available.

In another analysis we incorporate the priorities assigned by participants to the weights assigned to the 6 main characteristics as well as their interaction effects with their partner’s rating. Interaction effects in the linear regression setting refer to the product of the two variables, multiplying the score for a parameter with its associated importance, gives us a new value for each of the parameters. If people can correctly determine their priorities then there should be a direct correlation between the weighted scores given to a partner and the final decision (‘yes’ or ‘no’). Hence, our model is able to predict the most important factors in determining

the final decision of a participant. The logistic regression used in this analysis suggests that the interaction between partner attractiveness and priority of attractiveness, sincerity, and several other factors are significant in determining a participant's decision. Other parameters such as intelligence and being fun are not significant to the model. It is interesting to see that while in the previous analysis, partners being of the same race was not found to be significant, in this analysis the interaction between the variable for the importance of a partner being of the same race and the variable identifying whether the partners are of the same race was found to be highly significant to the model. This means that while, from an overall perspective, partners having the same race was not found to be significant, in cases where the participant said having a partner of the same race was important to them, it became significant.

5.2 Best Subset Selection Using Linear Models

We have also implemented the best subset selection model using exhaustive search, which studies all combinations of predictors for a given number of maximum predictors allowed in the model to determine the best possible combination. For this analysis we have used the previously studied variables such as scores given to a partner on the basis of attractiveness, sincerity, intelligence, being fun, ambitious and having shared interests, as well as other variables such as how much they liked the partner, the probability with which they think their partner will say yes to them, how happy they expect to be with the speed dating exercise, how often they go out (to see if extroverted behavior correlates with saying yes to more people), their goal behind attending the speed dating exercise, whether they have met the partner before and how often they go on dates. We have taken care to include only the variables that could directly impact outcome and have disregarded those that have a large number of missing values. Most of the variables with a large number of missing values were unrelated questions that we would not expect to impact the decision. The best subset in our model is selected using the Bayesian Information criterion (BIC) which is a criterion that uses both the maximum likelihood value as well as the number of parameters in the model to determine the best model overall. Thus, it takes into account how accurately the model fits the data while at the same time penalising models that have a large number of parameters. Hence, when determining the best model overall model, it gives us an accurate as well as simple model.

For a model containing a single predictor, the best one was found to be the score given for how much they liked the partner, this is clearly understandable as how much you like a person should have a very strong correlation with whether you say yes to them. For a model with two predictors, the best possible combination was found to be the score for how much they liked the partner and the score they gave for attractiveness of the partner. This is also clearly justifiable. Similarly, best possible subset combinations were predicted for subsets containing just one variable to those containing all thirteen variables. The best subset among all possible combinations of predictors and number of predictors was found to be the one containing nine variables, which were the scores given to the partner for attractiveness, sincerity, being fun, ambitious, having shared interests as well as how much they liked the partner, the probability with which they expect their partner to say yes to them, the score for how often they go out and the one for how happy they expect to be with the speed dating exercise.

While most of the predictors such as attractiveness, sincerity etc. are easy to justify, it is interesting to see that participants were more likely to say yes to a partner when they thought it was highly probable that their

partner would say yes to them. Also, participants who go out more often were also likely to say yes to more partners and those who expected to be happy with the speed dating exercise also said yes to more people. This indicates the possibility of a feedback loop, so participants who expect the speed dating exercise to go well for them are able to find more people that they are interested in.

5.3 Cross-Validation with Lasso Regularizer

In the logistic regression models described previously, we have used the entire dataset to create our model. While this ensures that we are using all the data available to us, thus creating a comprehensive model, it fails to tell us how well our model would test on data it has not previously seen and does not give us an indication of the variability of our model. The best way to check for that is to perform cross-validation, a technique that picks the best model on the basis of how well it performs when tested on the data it has not trained on.

As our goal is to create a model that is both accurate as well as easily interpretable, we have applied the technique of lasso regularization to our cross validation model. The lasso regularizer helps create a more sparse model, in which only the truly relevant parameters have non-zero coefficients, thus creating a simple solution.

For this model our input parameters remain the same as the ones used in the Best Subset model, i.e. the scores for the various attributes (attractiveness, sincerity etc.) and the other interesting variables (how much they liked the partner, how often they go out, how happy they expect to be etc.). The output remains the decision variable. For our cross-validation, since we are dealing with a dataset with 6716 entries, it was computationally impossible to implement the Leave One Out Cross-Validation (LOOCV), as that would require the model to be built on 6715 rows while being tested on only one outcome, and this itself would have to be repeated 6716 times. Instead we have chosen to implement a 20-fold cross-validation, which involves the entire dataset being divided into 20 individual sets. The model is then continuously built on 19 sets and tested on one, and this is repeated 20 times so that each set acts as the test set once, when the remaining 19 are being used to train the model.

In the glmnet package in the R coding language, the best value of λ (the coefficient of the lasso regularizer) is calculated by training the model on a subset of data with a number of different values of λ . Then, on the remaining testing data, the test error is calculated and the λ of the model that yields the minimizes the cross validation error is used. Hence, with all the above details we are able to implement our model, and get a good classification model.

5.4 K-Nearest-Neighbors (KNN)

In order to implement the K-Nearest-Neighbors (KNN) algorithm to predict a match between two people, the data must not have any missing values nor can it use categorical variables. Thus, some cleaning scripts were applied and some fields that had NA values were inputted to an “Other” category or the field was some form of weights or rankings, the values were imputed to be 0 instead of missing, just for the application of this algorithm. In addition, since there is relatively little data, a cross validation method is used to train and test the model on disjoint sections of data.

This method of reassigning or imputing may cause some bias in the data so we do not draw significant conclusions from this estimate. However, it is important that we do not skew the other results by imputing an average or

a most common value into some of the fields since it is possible that people who do not answer surveys have inherent traits in common so we want to try and bucket them together. In addition, it may be the case that the algorithm does not yield the same result for a match depending on the view (male or female as the participant) which has a skewed interpretation so we also try to fit whether or not the participant “likes” their partner. In addition, another thing to note is that since their partner has to have a different gender than the participant, currently we have gender included in the KNN model fit, but that should never be the same for the participant and the partner so better results may be achieved if that field was not included.

Initial trials of KNN showed that a few number of nearest neighbors to implement KNN is too flexible and we have a high misclassification rate. More tests are done to lower this number by increasing the number of nearest neighbors used and also using less fields in the algorithm.

6 Results and Key Findings

7 Next Steps

In the future we can utilize the information from follow-up surveys to the participant (a survey from the day after the speed dating event and another survey a few weeks after the first follow-up survey) as another measure of confirmation of the matches from the event, despite their low response rates and potential bias. However, when there are responses, they may give a lot of insight into what drives initial dates and calls.

In addition, so far when casually determining what predictors to use in the model, we have been using the full data set. However, in practice it is best to divide up the data that we have into disjoint training and test data so that we can train different models on the training data and then test the model on new data to get a more accurate estimate of test error. We have already developed a function to use the bootstrap method for resampling but that can also create biases and overfit the data. Another method that we will use for the linear or logistic models is the k-fold cross validation method to compute the error rate or number of misclassifications. This method begins by randomly dividing the data into k equally sized, disjoint partitions. Then each fold is taken in turn to be the test data set evaluated on the model that was calculated on the other $k - 1$ folds as the training set. In addition since there are only around 8,000 data points in total, the Leave-One-Out-Cross-Validation (LOOCV) method which is a special case of k-fold cross validation where k is equal to the total number of data points.